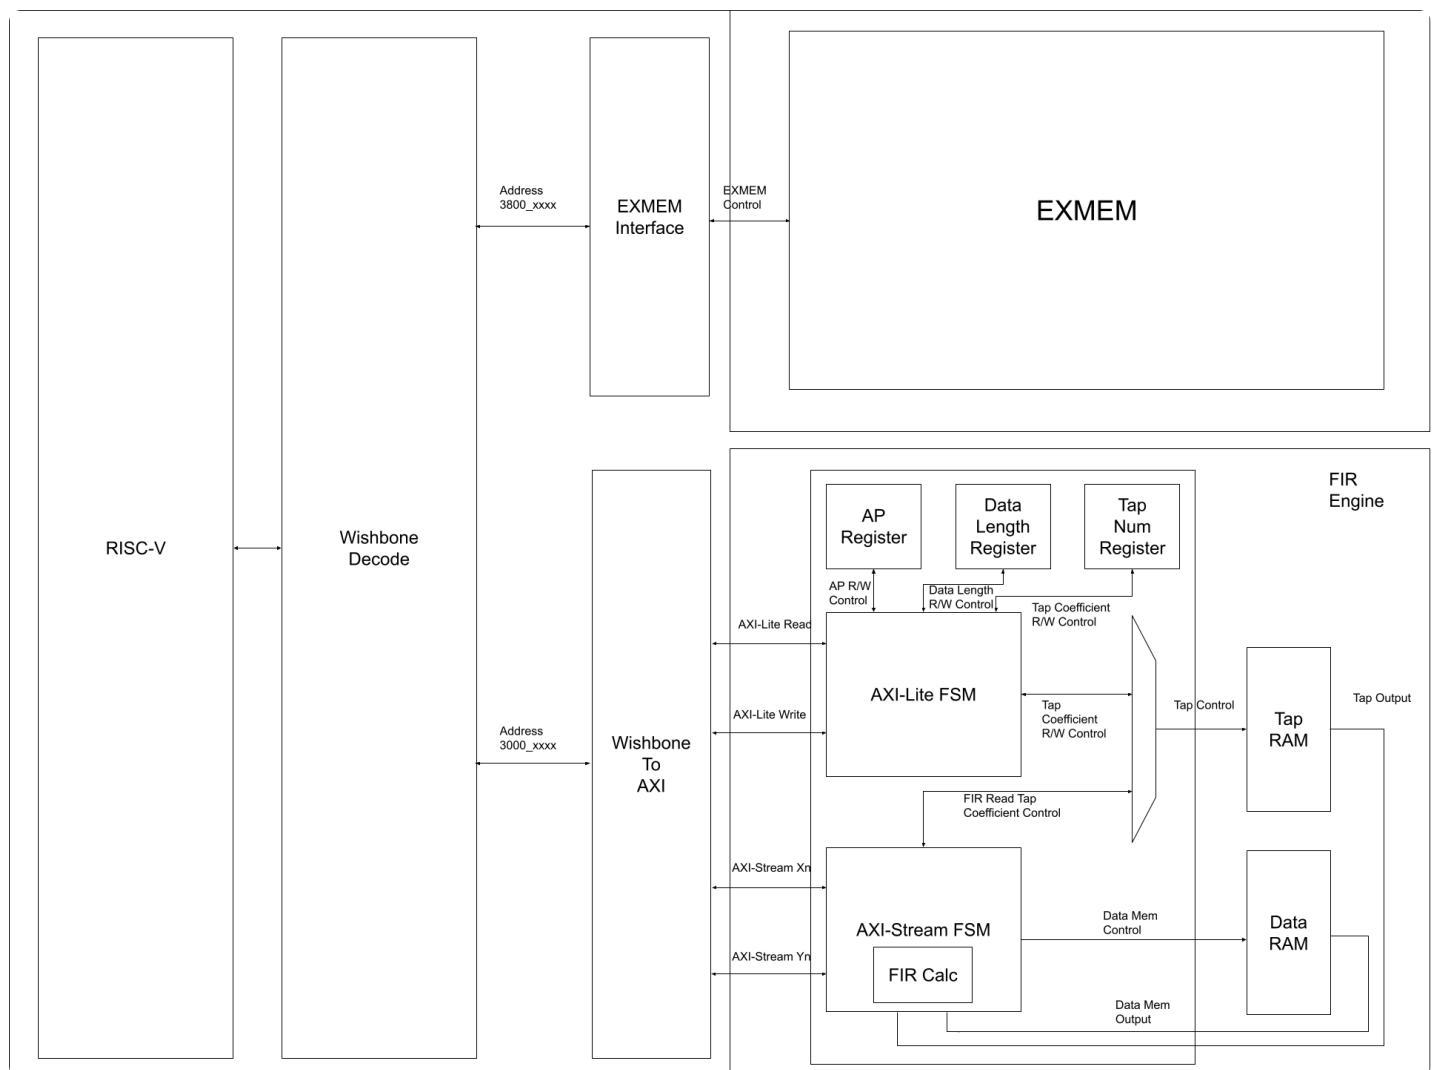


Lab 4-2 Caravel FIR

Student ID: 20726557

This lab involves integrating Lab 3-FIR and Lab 4-1 exmem-FIR into the Caravel user project area. A Wishbone to AXI interface is designed to communicate with FIR engine from Lab 3.

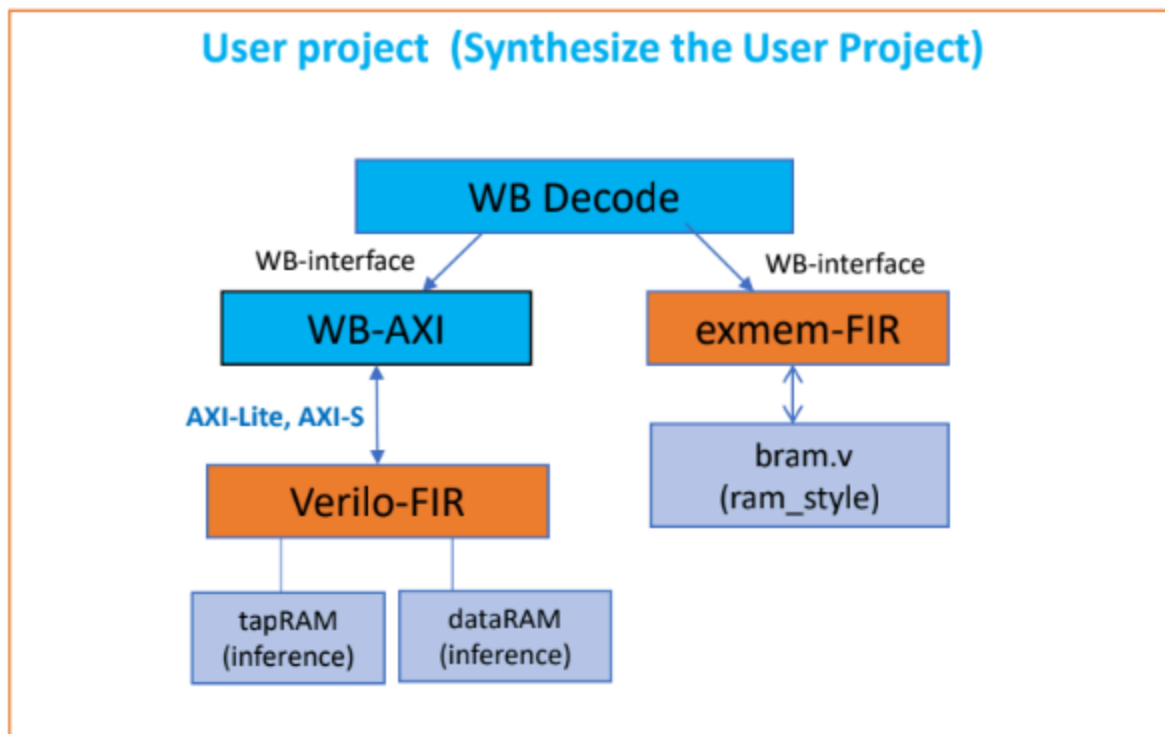
Design block diagram – datapath, control-path



The interface protocol between firmware, user project, and testbench

Firmware and User Project

The interface protocol between the firmware and the user project uses the Wishbone protocol to transmit data. The RISC-V CPU will execute the instructions in the firmware to transmit data through the Wishbone interface to the FIR engine and EXMEM. A Wishbone decoder will detect the access address of the transaction. The FIR engine is accessed by through the address of 0x3000_xxxx, while the EXMEM is accessed through the address of 0x3800_xxxx.



The firmware accesses the FIR engine in the User Project through a Wishbone-to-AXI conversion layer. The Wishbone-to-AXI layer detects the Wishbone request to access the FIR engine and converts the access into AXI-Lite or AXI-Stream format.

This lab uses the following addresses to access the FIR engine registers

- 0x00: AP Configuration Register (AXI-Lite)
- 0x10-13: Data Length Register (AXI-Lite)
- 0x14-18: Number of Taps (AXI-Lite)
- 0x40-43: X[n] input (AXI-Stream)

- 0x44-47: Y[n] input (AXI-Stream)
- 0x80-100: Tap parameters

The access of EXMEM is done directly through the Wishbone decoder to store and load the firmware code of the FIR engine. The Wishbone signals are converted into the bram control signals. The Wishbone receives an acknowledgement after a certain period of delay set by a parameter.

Firmware and Testbench

The Firmware and Testbench communicates through the mprj_io signal. The testbench sets mprj_io[23:16] as the checkbits to determine the state of the project. When the firmware loads the FIR engine, the checkbits will be set to 8'A5 and to notify the testbench to start checking for the FIR output and calculate the latency. After the FIR process has completed, the firmware sets the CPU to assert 8'h5A to the checkbits. The testbench will check the final Y[n] output with the a reference output and record the latency. This is repeated 3 times to and the total latency is recorded.

```


160      initial begin
161          wait(checkbits == 8'hA5);
162          $display("FIR Test 1 Started");
163
164          wait(checkbits == 8'h5A);
165          if(mprj_io[31:24] == 8'h76) begin
166              $display("FIR Test 1 Passed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
167              total_latency = total_latency + latency_count;
168          end else begin
169              $display("FIR Test 1 Failed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
170              $finish;
171          end
172
173          wait(checkbits == 8'hA5);
174          $display("FIR Test 2 Started");
175
176          wait(checkbits == 8'h5A);
177          if(mprj_io[31:24] == 8'h76) begin
178              $display("FIR Test 2 Passed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
179              total_latency = total_latency + latency_count;
180          end else begin
181              $display("FIR Test 2 Failed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
182              $finish;
183          end
184
185          wait(checkbits == 8'hA5);
186          $display("FIR Test 3 Started");
187
188          wait(checkbits == 8'h5A);
189          if(mprj_io[31:24] == 8'h76) begin
190              $display("FIR Test 3 Passed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
191              total_latency = total_latency + latency_count;
192              $display("Total Latency is %d", total_latency);
193          end else begin
194              $display("FIR Test 3 Failed, Final Y[n] is 0x%x, Latency is %d", mprj_io[31:24], latency_count);
195              $finish;
196          end
197          #10000;
198          $finish;
199      end

```

Waveform and analysis of the hardware/software behavior.

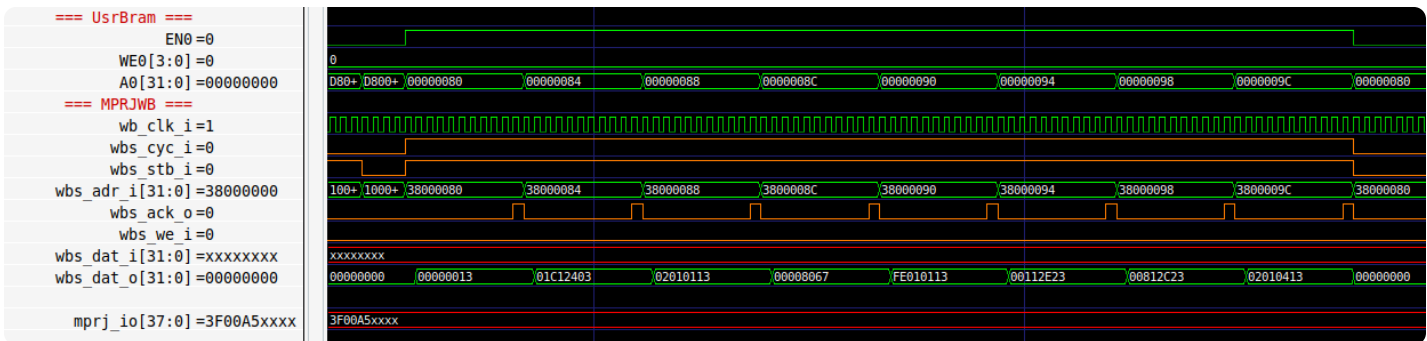
Software Behavior

Moving code from SPIFlash to EXMEM




The firmware writes the code to access the FIR engine from SPIFlash to the EXMEM. This is done through the Wishbone interface. After a certain delay set by a parameter, an acknowledge is sent back through the Wishbone interface to complete the write operation.

Reading code from EXMEM



When the FIR engine is enabled, the code is retrieved from the EXMEM. This is also done through the Wishbone interface and also follows the same delay to return the acknowledgement when the read is completed.

Initfir()




The FIR engine is initiated by passing the data length, number of taps, and the tap coefficients to the corresponding addresses. This is done in software by passing the data through using the Wishbone interface.

```

4 void __attribute__ ( ( section ( ".mprjram" ) ) ) initfir() {
5     //initial your fir
6     reg_fir_data_length = data_length;
7     reg_fir_tap_num = N;
8     for ( int i = 0; i < N; i++) {
9         reg_fir_coeff(i) = taps[i];
10    }
11 }
  
```


fir()



The software starts the FIR engine by setting AP configuration register to 0x01. The software transmits starts the X[n] and receives Y[n] through the Wishbone interface. The data length is 64 sets.

```
13 int* __attribute__ ( ( section ( ".mprjram" ) ) ) fir(){
14     initfir();
15     //write down your fir
16     reg_fir_control = 0x00000001;
17     for (int i = 0; i < data_length; i++) {
18         reg_fir_x = i;
19         outputsignal[i] = reg_fir_y;
20     }
21     return outputsignal;
22 }
```

Communication with Testbench

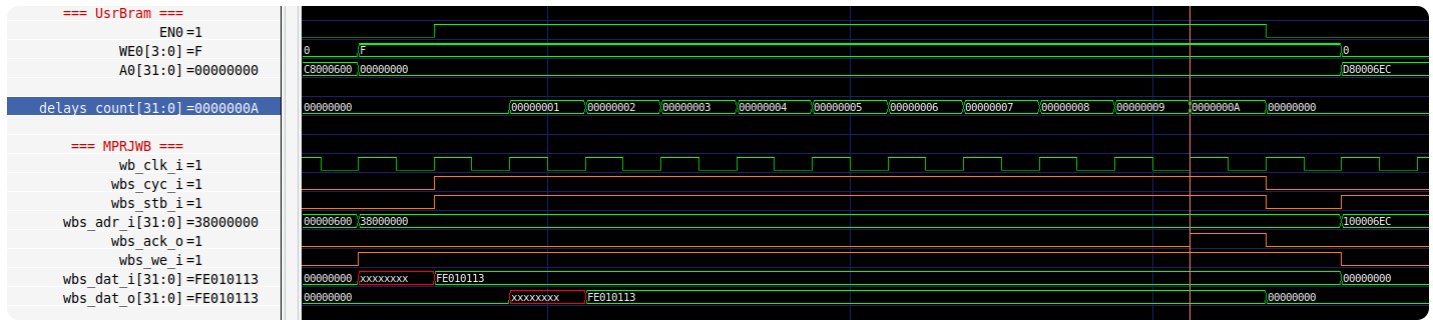


The software communicates with the verilog testbench through the mprj_io. When the checkbits is set to 8'hA5, the testbench will start the latency counter. When the FIR engine completes, the software sets the checkbits to 8'5A, allowing the testbench to determine the total latency.

```
119     for (int i = 0; i < 3; i++) {
120         // Flag start of the test
121         reg_mprj_data1 = 0x00A50000;
122
123         int* tmp = fir();
124         reg_mprj_data1 = (*(tmp + 63) << 24) + 0x005A0000;
125
126         //print("\n");
127         //print("Monitor: Test 1 Passed\n\n");    // Makes simulation very long!
128         //reg_mprj_data1 = 0x005A0000;
129     }
130 }
```

Hardware Behavior

Wishbone-to-EXMEM Delay Interface

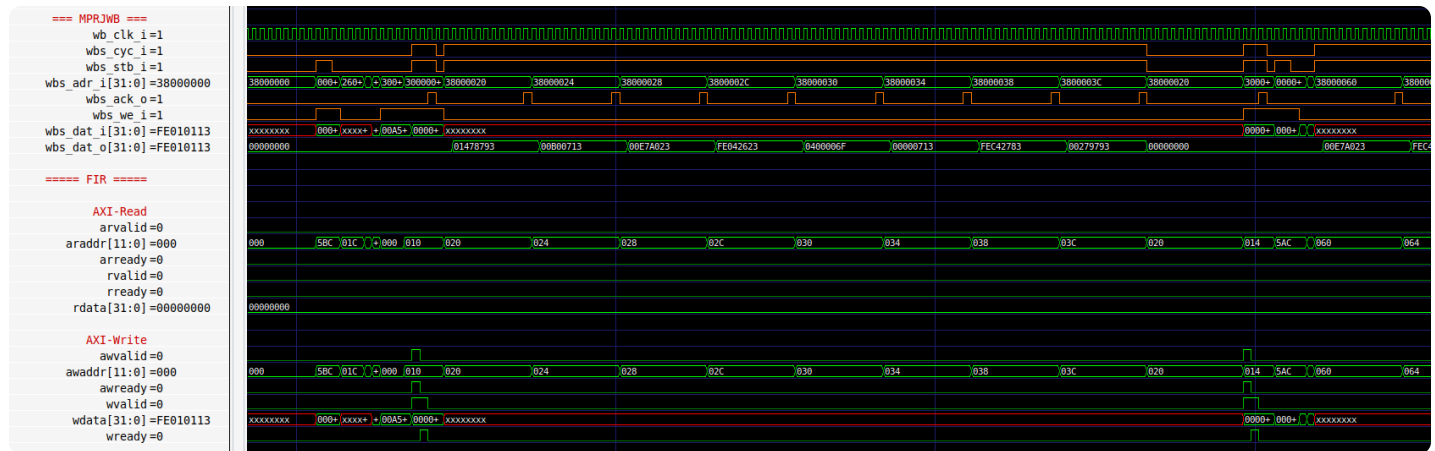


The Wishbone-to-EXMEM Delay Interface is in charge of the read/write operation between the Firmware and the EXMEM. The interface converts the Wishbone signals into BRAM control signals such that the data can be read/written. The wbs_ack_o acknowledgement signal is asserted after a certain delay, determined using the DELAYS parameter, through the use of a delay counter.

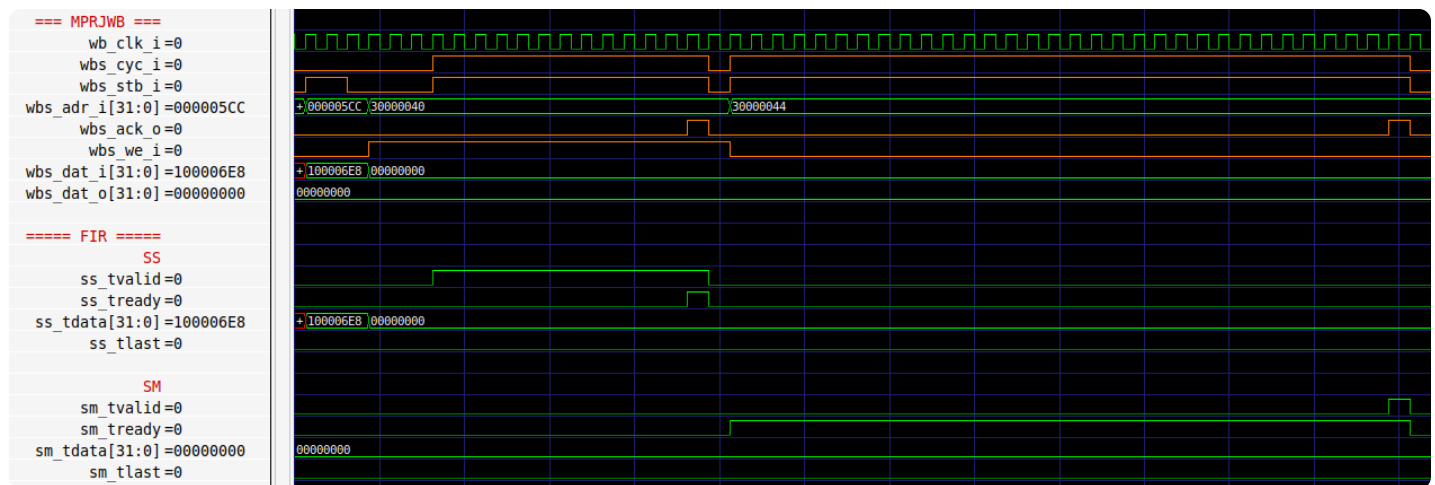
```
283 // EXMEM
284 assign exmem_en      = wbs_stb_i & wbs_cyc_i & (wbs_adr_i[31:24] == 8'h38);
285 assign exmem_we[3:0] = {4{wbs_we_i}} & wbs_sel_i[3:0];
286 assign exmem_adr[31:0] = wbs_adr_i[31:0] - 32'h38000000;
287
288 always @(posedge wb_clk_i) begin
289     if (wb_rst_i) begin
290         delays_count[31:0] <= 32'b0;
291     end else begin
292         if (exmem_en) begin
293             if (delays_count[31:0] != DELAYS) begin
294                 delays_count[31:0] <= delays_count[31:0] + 32'b1;
295             end else begin
296                 delays_count[31:0] <= 32'b0;
297             end
298         end
299     end
300 end
301
302 bram user_bram (
303     .CLK(wb_clk_i),
304     .WE0(exmem_we[3:0]),
305     .EN0(exmem_en),
306     .Di0(wbs_dat_i[31:0]),
307     .Do0(exmem_dat_o[31:0]),
308     .A0 (exmem_adr[31:0])
309 );
```

Wishbone-to-AXI Interface

AXI-Lite



AXI-Stream




To communicate with the FIR engine, the Wishbone signals from the CPU is converted to the AXI protocol. To access the AP Configuration Register, Data Length, Tap Number registers, and Tap Coefficient RAM, the Wishbone signals are converted in AXI-Lite Protocol. The Read and Write transactions of the FIR engine is converted from Wishbone to AXI-Stream Protocol.


```

225 assign fir_en = wbs_stb_i & wbs_cyc_i & (wbs_adr_i[31:24] == 8'h30);
226
227 // AXI-Lite
228 assign axilite_en = fir_en & ~((wbs_adr_i[6:0] >= 7'h40) & (wbs_adr_i[6:0] <= 7'h47));
229 assign axilite_adr[31:0] = wbs_adr_i[31:0] - 32'h30000000;
230
231 always @(posedge wb_clk_i or posedge wb_rst_i) begin
232     if (wb_rst_i) begin
233         axilite_awready <= 1'b0;
234         axilite_wready <= 1'b0;
235         axilite_arready <= 1'b0;
236     end else begin
237         if (wbs_ack_o) begin
238             axilite_awready <= 1'b0;
239             axilite_wready <= 1'b0;
240             axilite_arready <= 1'b0;
241         end else begin
242             axilite_awready <= awready | axilite_awready;
243             axilite_wready <= wready | axilite_wready;
244             axilite_arready <= arready | axilite_arready;
245         end
246     end
247 end
248
249 assign awvalid = axilite_en & wbs_we_i & ~axilite_awready;
250 assign awaddr[(pADDR_WIDTH-1):0] = axilite_adr[(pADDR_WIDTH-1):0];
251
252 assign wvalid = axilite_en & wbs_we_i & ~axilite_wready;
253 assign wdata[(pDATA_WIDTH-1):0] = wbs_dat_i[(pDATA_WIDTH-1):0];
254
255 assign arvalid = axilite_en & ~wbs_we_i & ~axilite_arready;
256 assign araddr[(pADDR_WIDTH-1):0] = axilite_adr[(pADDR_WIDTH-1):0];
257
258 assign rready = axilite_en & ~wbs_we_i;
259
260 // AXI-Stream
261 assign axistream_x = fir_en & (wbs_adr_i[6:0] == 7'h40);
262 assign axistream_y = fir_en & (wbs_adr_i[6:0] == 7'h44);
263
264 always @(posedge wb_clk_i or posedge wb_rst_i) begin
265     if (wb_rst_i) begin
266         axistream_count[31:0] <= 32'b0;
267     end else begin
268         if (axistream_x & wbs_ack_o) begin
269             axistream_count[31:0] <= axistream_count[31:0] + 32'b1;
270         end else begin
271             axistream_count[31:0] <= axistream_count[31:0];
272         end
273     end
274 end
275
276 assign ss_tvalid = axistream_x & wbs_we_i;
277 assign ss_tdata[(pDATA_WIDTH-1):0] = wbs_dat_i[(pDATA_WIDTH-1):0];
278 assign ss_tlast = ss_tvalid & (axistream_count[31:0] == pDATA_LENGTH - 1);
279
280 assign sm_tready = axistream_y & ~wbs_we_i;
281


```

FIR Engine



The FIR Engine receives the AXI-Lite and AXI-Stream transactions to compute the FIR outputs. When `ap_start` is set to 1, the FIR engine starts the computation by reading `X[n]` and output `Y[n]` through the AXI-Stream.

What is the FIR engine theoretical throughput, i.e. data rate? Measured throughput?



The theoretical delay per $Y[n]$ output is 33 cycles. With a data length of 64, the theoretical delay is 2112 cycles.

Therefore, the thoretical throughput per 64 data sets is the following:

Theoretical throughput: $1/2112 = 0.00047348484$ data per cycle

However, the actual measured delay for a data length of 64 is larger than expected.

```
ubuntu@ubuntu2004:~/Desktop/EESM6000C/caravel-soc_fpga-lab/lab-caravel_fir/testbench/counter_la_fir$ source run_sim
Reading counter_la_fir.hex
counter_la_fir.hex loaded into memory
Memory 5 bytes = 0x6f 0x00 0x00 0x0b 0x13
VCD info: dumpfile counter_la_fir.vcd opened for output.
FIR Test 1 Started
FIR Test 1 Passed, Final Y[n] is 0x76, Latency is      40006
FIR Test 2 Started
FIR Test 2 Passed, Final Y[n] is 0x76, Latency is      40006
FIR Test 3 Started
FIR Test 3 Passed, Final Y[n] is 0x76, Latency is      39995
Total Latency is      120007
```

Measured throughput: $1/40006 = 0.00002499625$ data per cycle

The cause of the increase in delay from theoretical is due to the time between each $X[n]$ input. The FIR engine stalls and waits until the next $X[n]$ input is received through the AXI-Stream.

What is the latency for firmware to feed data?

The latency for the firmware to feed data can be determined through the cycles between each $X[n]$ input from AXI-Stream.



Latency: $(1586537500 - 1574937500) \text{ ps} / (25000) \text{ ps} = 464 \text{ cycles}$

What techniques are used to improve the throughput?

There are several techniques that can be done to improve the throughput. In terms of hardware, the FIR engine can be optimized.

The FIR engine can only receive $X[n]$ and send $Y[n]$ one at a time. If there is no data sent from the RISC-V CPU to the FIR engine, the FIR engine will stall, affecting the throughput. To improve this, a buffer can be added to store $X[n]$ inputs and $Y[n]$ output, and will not stall the FIR engine even if the SS/SM Bus tvalid is not asserted.

Another optimization done in the FIR engine is parallelism. In the current design, multiply and sum actions are done in separate cycles. The FIR engine can take advantage of the parallelism by using the Multiplier and Adder at the same cycles. When the previous dataset is using the Adder, the next dataset can use the Multiplier. This will reduce the FIR calculation delay and improve the throughput.

In terms of software, the throughput can be improved as well. There is a large latency for the firmware to feed the data. To improve this, the $X[n]$ input can be decoupled from the $Y[n]$ output, and allow $X[n]$ to input into the user project area before receiving the $Y[n]$ output. This will improve the delay between each data input and improve the throughput.

Any other insights?

Synthesis

Utilization Report

31		Used	Fixed	Prohibited	Available	Util%
32	Site Type					
33						
34	Slice LUTs*	4802	0	0	53200	9.03
35	LUT as Logic	4684	0	0	53200	8.80
36	LUT as Memory	118	0	0	17400	0.68
37	LUT as Distributed RAM	80	0			
38	LUT as Shift Register	38	0			
39	Slice Registers	4205	0	0	106400	3.95
40	Register as Flip Flop	4130	0	0	106400	3.88
41	Register as Latch	75	0	0	106400	0.07
42	F7 Muxes	169	0	0	26600	0.64
43	F8 Muxes	47	0	0	13300	0.35
44						

Site Type	Used	Fixed	Prohibited	Available	Util%
Block RAM Tile	5	0	0	140	3.57
RAMB36/FIFO*	2	0	0	140	1.43
RAMB36E1 only	2				
RAMB18	6	0	0	280	2.14
RAMB18E1 only	6				

```
85 +-----+-----+-----+-----+-----+-----+
86 | Site Type | Used | Fixed | Prohibited | Available | Util% |
87 +-----+-----+-----+-----+-----+-----+
88 | DSPs      |    0 |     0 |           0 |         220 |   0.00 |
89 +-----+-----+-----+-----+-----+-----+
90
```

162	-----											
163	Design Timing Summary											
164	-----											
165												
166												
167	WNS(ns)	TNS(ns)	TNS Failing Endpoints	TNS Total Endpoints	WHS(ns)	THS(ns)	THS Failing Endpoints	THS Total Endpoints	WPWS(ns)	TPWS(ns)	TPWS Failing Endpoints	TPWS Total
168	Endpoints											
169	38.889	0.000	0	13395	0.032	0.000	0	13395	48.750	0.000		
170	0	5418										
171												
172	All user specified timing constraints are met.											

```

230 Max Delay Paths
231 -----
232 Slack (MET) : 38.889ns (required time - arrival time)
233 Source: design_1_i/processing_system7_0/inst/PS7_i/FCLKCLK[0]
234 (clock source 'clk_fpga_0' {rise@0.000ns fall@50.000ns period=100.000ns})
235 Destination: design_1_i/caravel_0/inst/housekeeping/wb_dat_o_reg[7]/D
236 (rising edge-triggered cell FDSE clocked by clk_fpga_0 {rise@0.000ns fall@50.000ns period=100.000ns})
237 Path Group: clk_fpga_0
238 Path Type: Setup (Max at Slow Process Corner)
239 Requirement: 50.000ns (clk_fpga_0 rise@100.000ns - clk_fpga_0 fall@50.000ns)
240 Data Path Delay: 12.097ns (logic 1.231ns (10.176%) route 10.866ns (89.824%))
241 Logic Levels: 8 (BUFG=1 LUT3=1 LUT4=1 LUT6=4 MUXF7=1)
242 Clock Path Skew: 2.716ns (DCD - SCD + CPR)
243 Destination Clock Delay (DCD): 2.716ns = ( 102.716 - 100.000 )
244 Source Clock Delay (SCD): 0.000ns = ( 50.000 - 50.000 )
245 Clock Pessimism Removal (CPR): 0.000ns
246 Clock Uncertainty: 1.500ns ((TSJ^2 + TIJ^2)^1/2 + DJ) / 2 + PE
247 Total System Jitter (TSJ): 0.071ns
248 Total Input Jitter (TIJ): 3.000ns
249 Discrete Jitter (DJ): 0.000ns
250 Phase Error (PE): 0.000ns
251
252 Location Delay type Incr(ns) Path(ns) Netlist Resource(s)
253 -----
254 (clock clk_fpga_0 fall edge)
255
256 PS7_X0Y0 PS7 0.000 50.000 f design_1_i/processing_system7_0/inst/PS7_i/FCLKCLK[0]
257 net (fo=1, routed) 1.193 51.193 design_1_i/processing_system7_0/inst/FCLK_CLK_unbuffered[0]
258 BUFGCTRL_X0Y21 BUFG (Prop_bufg_I_0) 0.101 51.294 f design_1_i/processing_system7_0/inst/buffer_fclk_clk_0.FCLK_CLK_0_BUFG/0
259 net (fo=5421, routed) 2.084 53.378 design_1_i/caravel_0/inst/gpio_control_in_1[7]/clock
260 SLICE_X50Y92 LUT6 (Prop_lut6_I3_0) 0.124 53.502 f design_1_i/caravel_0/inst/gpio_control_in_1[7]/mprj_o[15]_INST_0/0
261 net (fo=2, routed) 1.861 55.363 design_1_i/caravel_ps_0/inst/control_s_axi_U/mprj_in[15]_INST_0/0
262 SLICE_X45Y91 LUT3 (Prop_lut3_I0_0) 0.124 55.487 f design_1_i/caravel_ps_0/inst/control_s_axi_U/mprj_in[15]_INST_0/0
263 net (fo=1, routed) 1.057 56.543 design_1_i/caravel_0/inst/housekeeping/hkspi/mprj_i[15]
264 SLICE_X45Y91 LUT6 (Prop_lut6_I0_0) 0.124 56.667 f design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[15]_i_22/0
265 net (fo=1, routed) 1.067 57.734 design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[15]_i_22_n_0
266 SLICE_X60Y89 LUT6 (Prop_lut6_I0_0) 0.124 57.858 f design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[15]_i_14/0
267 net (fo=1, routed) 0.000 57.858 design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[15]_i_14_n_0
268 SLICE_X60Y89 MUXF7 (Prop_muxf7_I1_0) 0.217 58.075 f design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o_reg[15]_i_7/0
269 net (fo=1, routed) 0.851 58.926 design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o_reg[15]_i_7_n_0
270 SLICE_X57Y87 LUT6 (Prop_lut6_I4_0) 0.299 59.225 f design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[15]_i_3/0
271 net (fo=5, routed) 2.048 61.273 design_1_i/caravel_0/inst/housekeeping/hkspi/wbdat_busy_reg_29
272 SLICE_X53Y52 LUT4 (Prop_lut4_I3_0) 0.118 61.391 f design_1_i/caravel_0/inst/housekeeping/hkspi/wb_dat_o[7]_i_2/0
273 net (fo=1, routed) 0.705 62.097 design_1_i/caravel_0/inst/housekeeping/hkspi_n_90
274 SLICE_X54Y52 FDSE f design_1_i/caravel_0/inst/housekeeping/wb_dat_o_reg[7]/D
275 -----
276
277 (clock clk_fpga_0 rise edge)
278
279 PS7_X0Y0 PS7 0.000 100.000 r design_1_i/processing_system7_0/inst/PS7_i/FCLKCLK[0]
280 net (fo=1, routed) 1.088 101.088 design_1_i/processing_system7_0/inst/FCLK_CLK_unbuffered[0]
281 BUFGCTRL_X0Y21 BUFG (Prop_bufg_I_0) 0.091 101.179 r design_1_i/processing_system7_0/inst/buffer_fclk_clk_0.FCLK_CLK_0_BUFG/0
282 net (fo=5421, routed) 1.537 102.716 design_1_i/caravel_0/inst/housekeeping/clock
283 SLICE_X54Y52 FDSE r design_1_i/caravel_0/inst/housekeeping/wb_dat_o_reg[7]/C
284 clock pessimism 0.000 102.716
285 clock uncertainty -1.500 101.215
286 SLICE_X54Y52 FDSE (Setup_fdse_C_D) -0.230 100.985 design_1_i/caravel_0/inst/housekeeping/wb_dat_o_reg[7]
287 -----
288 required time 100.985
289 arrival time -62.097
290 -----
291 slack 38.889

```

GitHub

<https://github.com/AnthonyGithub/EESM6000C-Lab-4/tree/main/Lab 4-2>

(<https://github.com/AnthonyGithub/EESM6000C-Lab-4/tree/main/Lab%204-2>)